

# MicroWatt Microarchitecture

H. Peter Hofstee

(uArch overview largely based on a slide deck by Paul Mackerras)

---

This slide deck is not IBM confidential

# Contents

- MicroWatt Microarchitecture
- Adding new Instructions to MicroWatt

## Goals for today:

- Intro to pipelining
- Learn more about Power instructions/formats & instruction recode/decode
- Intro to branch prediction
- Intro to cache design
- Show some vhd

## Next Tue (IST)

- Front end topics: Branch Prediction, Register Renaming, Dependency Checking

## Next Th (IST)

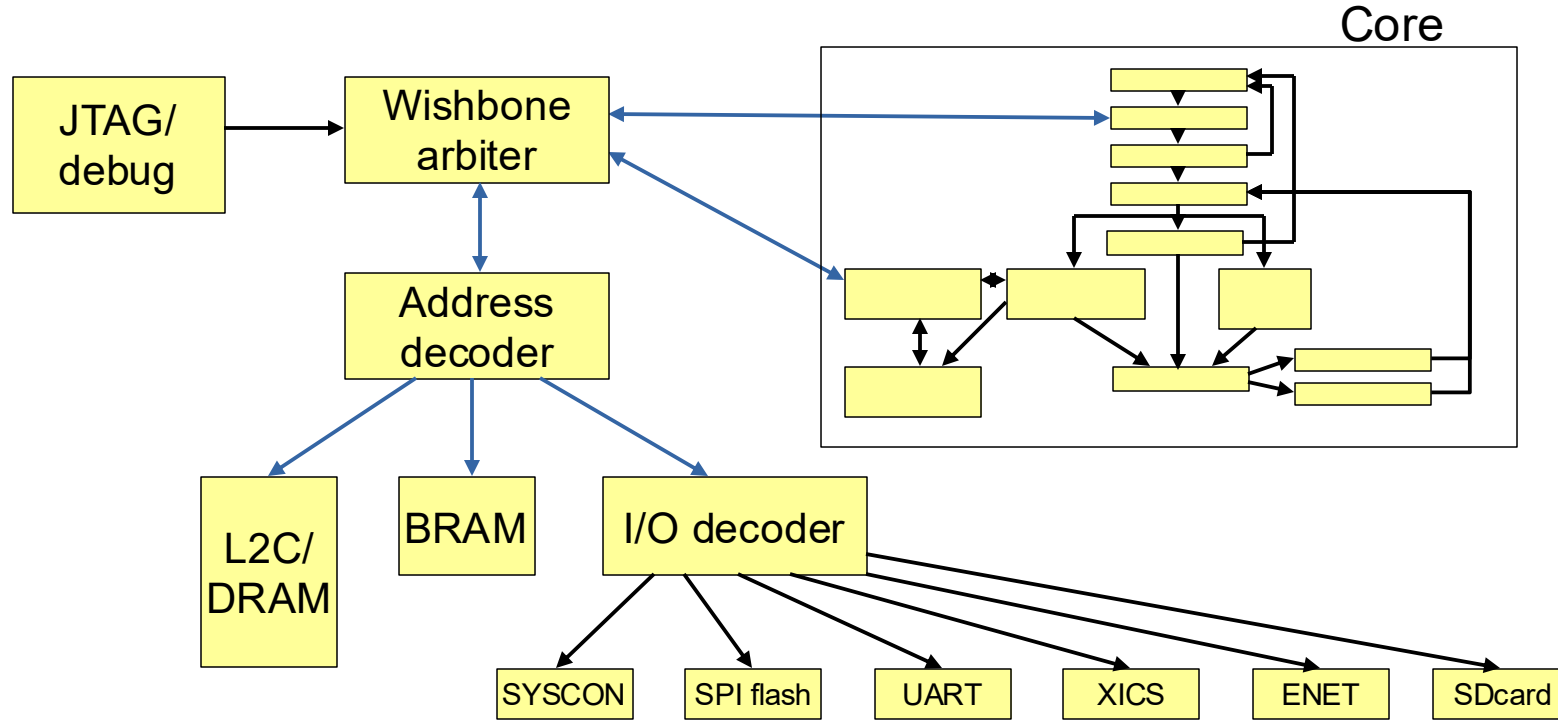
- Back end topics: Load-Store Unit, Instruction Writeback and Completion

# Introduction



- Microwatt started as demo/proof-of-concept for announcement of Power ISA being made open (August 2019).
  - Core structure is relatively simple
- Code hosted on github, updated through pull requests
  - <https://github.com/antonblanchard/microwatt>
  - Use automated testing to catch bugs early
- Aims to be a compliant PowerISA v3.1C implementation
  - Includes support for prefixed instructions (2-word instructions)
- Can run Linux
- Targets simulation and synthesis for FPGAs
  - E.g. Artix-7, see [https://github.com/hofstee-hp/MicroWatt\\_on\\_Arty\\_A7-100T](https://github.com/hofstee-hp/MicroWatt_on_Arty_A7-100T)
- Wishbone interface to memory (and memory-mapped peripherals)
- Peripherals from Litex project



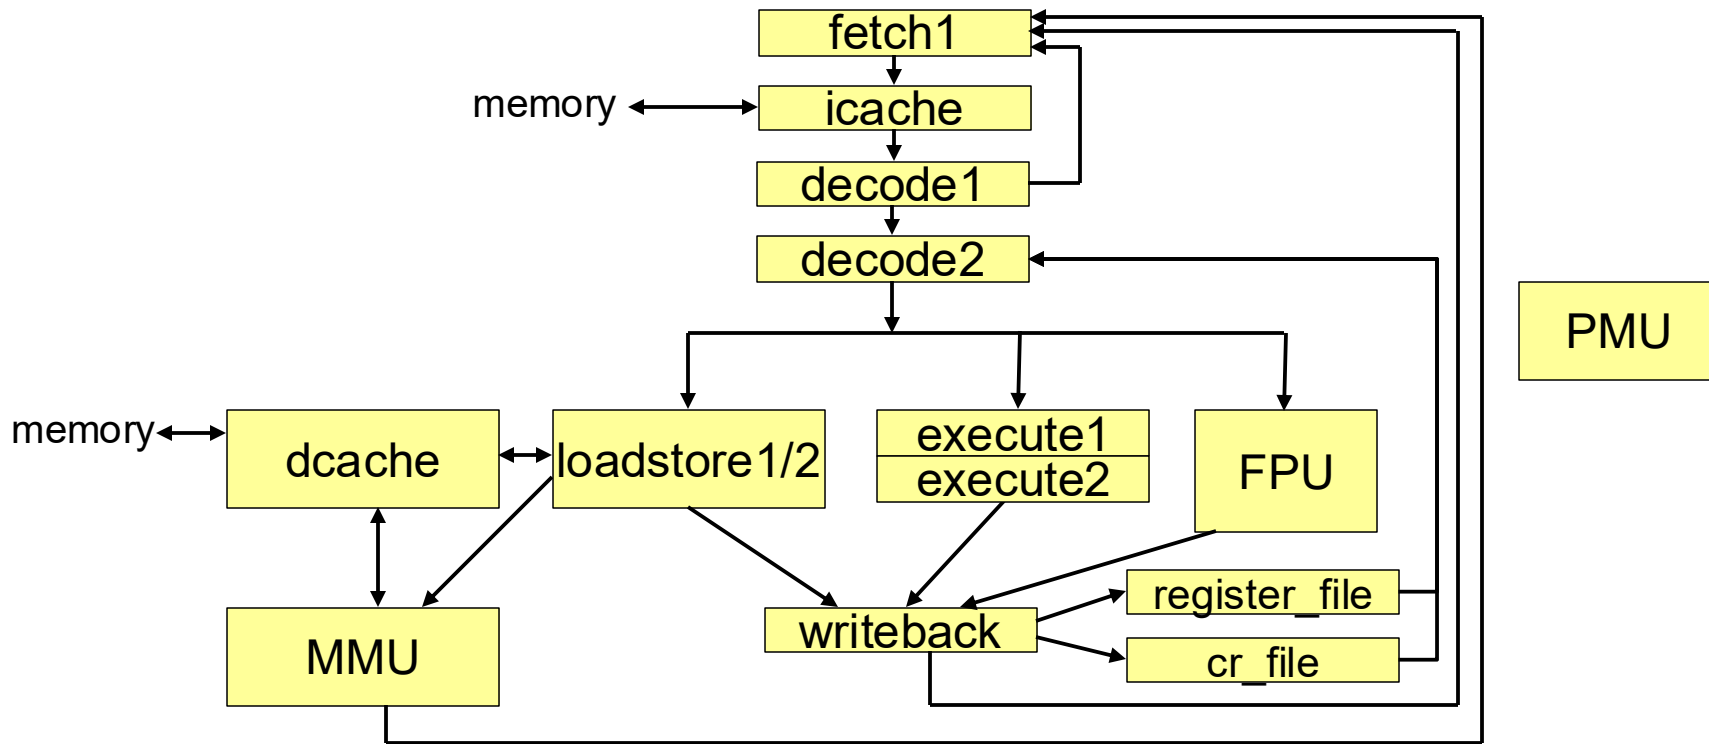
# System overview



# Arty A7-100T



# Pipeline Overview



# Fetch1 and icache

- Fetch1 generates a sequence of effective and corresponding real addresses
  - Increment by 4 each cycle unless redirected by decode1 or writeback
    - Branches, interrupts, rfid, isync cause redirects
  - Contains instruction TLB and tiny 2-entry cache of TLB
  - Contains (optional) Branch Target Cache (BTC)
    - Stores instruction address, target address and taken indication for executed direct branches
- Icache reads and caches instructions from memory
  - Use of real address for index and tag avoids aliasing problems
  - Instructions are predecoded into 36-bit form on icache refill
    - 6-bit primary opcode gets replaced by 10-bit instruction index determined from primary and extended opcodes
  - No support currently from fetching from cache-disabled pages of memory
  - Icache snoops writes to memory and invalidates corresponding lines



# Instruction Opcode Recoding (predecode.vhdl)

- OpenPOWER has MANY instruction formats
  - pp. 11-16 of v3.1c ISA document
- All have a 6 bit primary opcode (PO)
  - Extended opcodes (XO) range from 0 to 11 additional bits
  - Sometimes other modifier bits also effectively change the operation
- Many POWER processors recode instruction bits before storing in the iCache for easier decode in the rest of the pipeline
  - Not only opcodes (as in microwatt) some also regularize register operand locations etc.

architecture behaviour of predecoder is

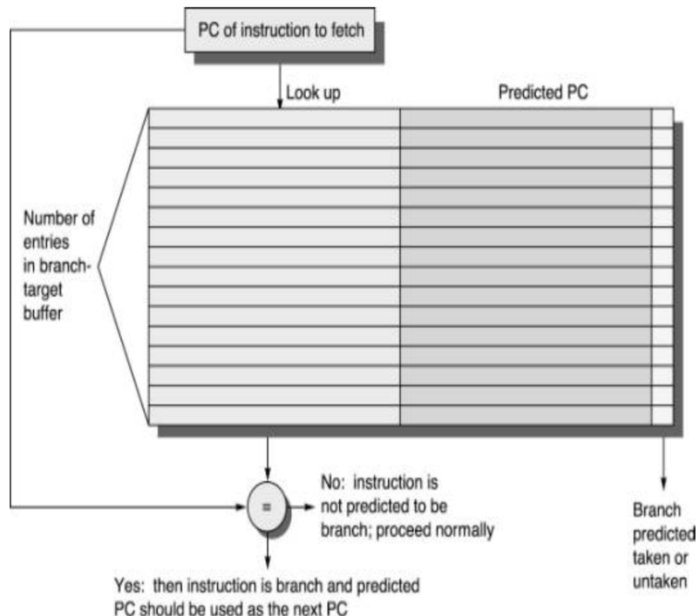
```
type predecoder_rom_t is array(0 to 2047) of insn_code;
```

```
constant major_predecode_rom : predecoder_rom_t := (  
    2#001100_00000# to 2#001100_11111# => INSN_addic,  
    2#001101_00000# to 2#001101_11111# => INSN_addic_dot,  
    2#001110_00000# to 2#001110_11111# => INSN_addi,  
    2#001111_00000# to 2#001111_11111# => INSN_addis,  
    2#010011_00100# to 2#010011_00101# => INSN_addpcis,  
    2#011100_00000# to 2#011100_11111# => INSN_andi_dot,  
    2#011101_00000# to 2#011101_11111# => INSN_andis_dot,  
    2#000000_00000# to 2#000000_11111# => INSN_attn,  
    2#010010_00000# to 2#010010_00001# => INSN_brel,  
    2#010010_00010# to 2#010010_00011# => INSN_babs,  
    2#010010_00100# to 2#010010_00101# => INSN_brel,  
    2#010010_00110# to 2#010010_00111# => INSN_babs,  
    2#010010_01000# to 2#010010_01001# => INSN_brel,  
    2#010010_01010# to 2#010010_01011# => INSN_babs,  
    2#010010_01100# to 2#010010_01101# => INSN_brel,  
    2#010010_01110# to 2#010010_01111# => INSN_babs,  
    2#010010_10000# to 2#010010_10001# => INSN_brel,  
    2#010010_10010# to 2#010010_10011# => INSN_babs,  
    2#010010_10100# to 2#010010_10101# => INSN_brel,  
    2#010010_10110# to 2#010010_10111# => INSN_babs,  
    2#010010_11000# to 2#010010_11001# => INSN_brel,  
    2#010010_11010# to 2#010010_11011# => INSN_babs,  
    2#010010_11100# to 2#010010_11101# => INSN_brel,  
    2#010010_11110# to 2#010010_11111# => INSN_babs,
```





# Branch Target Address Cache (Fetch1)




```
-- If there is a valid entry in the BTC which corresponds to the next instruction,
-- use that to predict the address of the instruction after that.
-- (w_in.redirect = '0' and d_in.redirect = '0' and r_int.tlbstall = '0')
-- implies v.nia = r_int.next_nia.
-- r_int.rd_is_niap4 implies r_int.next_nia is the address used to read the BTC.
if v.req = '1' and w_in.redirect = '0' and d_in.redirect = '0' and r_int.tlbstall = '0' and
    btc_rd_valid = '1' and r_int.rd_is_niap4 = '1' and
    btc_rd_data(BTC_WIDTH - 2) = r.virt_mode and
    btc_rd_data(BTC_WIDTH - 3 downto BTC_TARGET_BITS)
        = r_int.next_nia(BTC_TAG_BITS + BTC_ADDR_BITS + 1 downto BTC_ADDR_BITS + 2) then
    v.predicted := btc_rd_data(BTC_WIDTH - 1);
    v.pred_ntaken := not btc_rd_data(BTC_WIDTH - 1);
    if btc_rd_data(BTC_WIDTH - 1) = '1' then
        v_int.next_nia := btc_rd_data(BTC_TARGET_BITS - 1 downto 0) & "00";
        v_int.rd_is_niap4 := '0';
    end if;
end if;
```




# Virtually vs. Physically Addressed Cache

## Physical Cache



## Virtual Cache



## Physical Cache

**Slow:** Must do a TLB lookup *before* accessing the cache

## Virtual Cache

**Fast:** TLB lookups *only* when we miss in the cache

**Q:** Can you have two programs share a virtual cache?

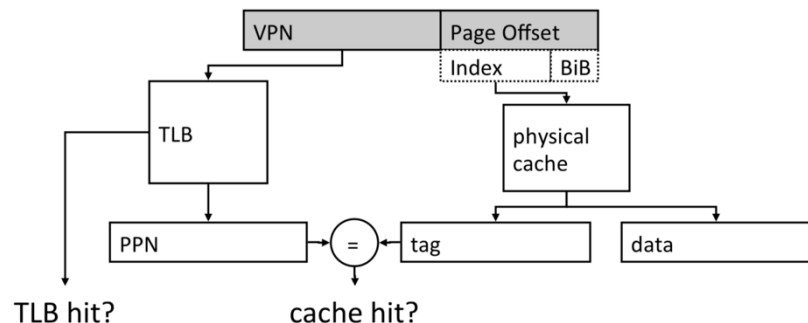
- ☐ Yes. The TLB provides protection.
- ☐ No. Each needs its own TLB.
- ☐ No. The cache is virtual so there is no way to provide protection

**A:** No.

A virtual cache stores data by the virtual program address. There is no translation, so VM-based protection can't keep applications apart!

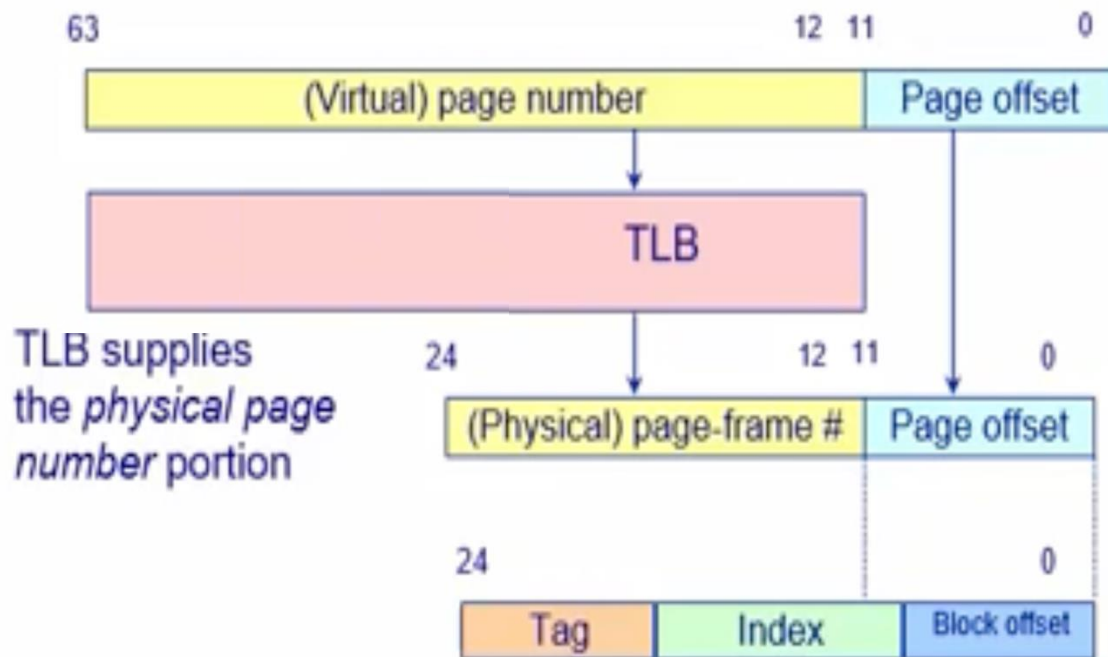
# Basic Cache Types

- Physically Indexed, Physically Tagged (PIPT)
  - Straightforward, but slow, because address translation has to be done first
  - Microwatt iCache design keeps last two TLB entries
- Virtually Indexed, Virtually Tagged
  - Fast (no translation required), but ...
  - ....need to clear on context switch (because of security and aliasing issues)
- Virtually Indexed, Physically Tagged (VIPT)
  - To do translation in parallel w. cach access each set is limited to the page size

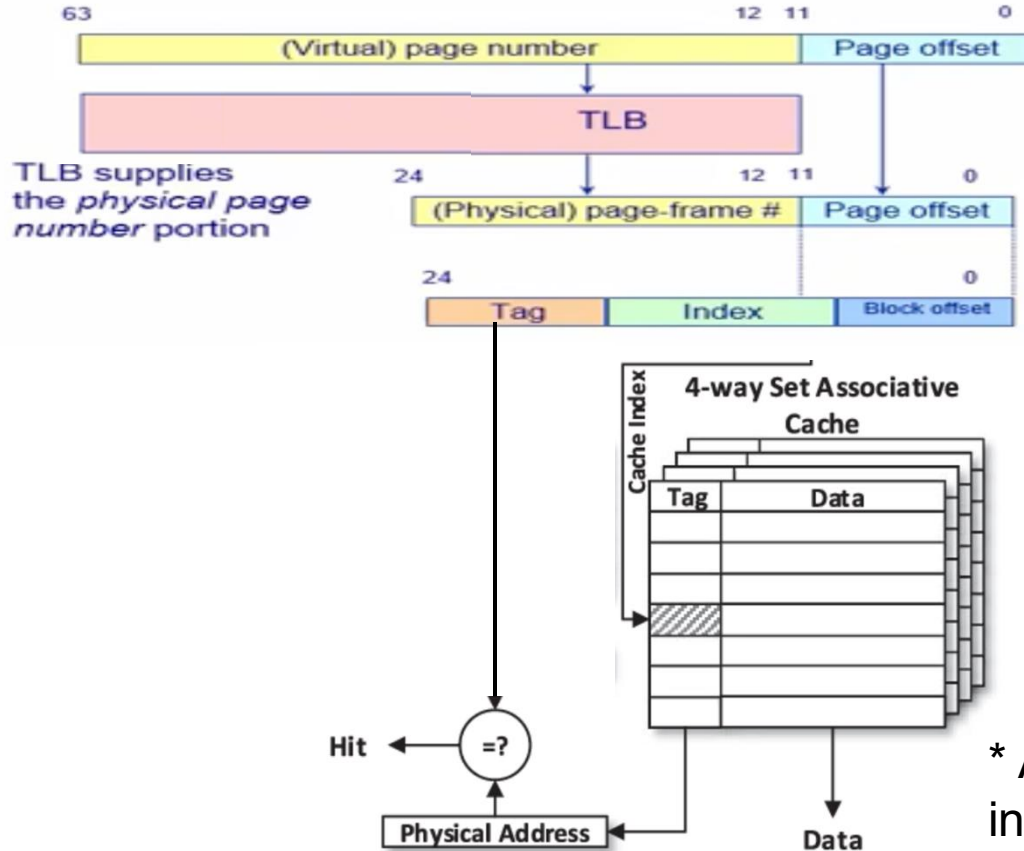


<http://thebeardsage.com/virtual-memory-cache-tlb-interaction/>

# Virtual vs. Real/Physical Address



# Cache Access



\* Microwatt uses a 2-entry cache of the TLB to avoid having to access the TLB for every iFetch.

\* At very least also need valid indication per entry (also replacement-related bits).

# Decode1 and decode2

- Decode1 looks up a decode ROM using instruction index
  - Decode ROM gives control signals used to control operand decoding, ALU function, result selection, etc.
  - Does static branch prediction for direct branches
  - Computes register file addresses (GPRs and FPRs) for up to 3 register operands
  - Decodes SPR (special purpose register) numbers
- Decode2 does instruction scheduling and dispatch
  - Calculates instruction dependencies and stalls until operands are available
  - Forwards operands from previous instruction results (where available)
  - Does immediate operand selection and formatting
    - Operand can come from instruction field, or be the address of the instruction
  - Computes some control signals for execute stage



# Decode1 decoder ROM

```
type decoder_rom_t is array(insn_code) of decode_rom_t;
```

```
constant decode_rom : decoder_rom_t := (
```

|                 | unit | fac    | internal | in1              | in2         | const             | in3         | out  | CR   | CR   | inv  | inv   | cry  | cry   | ldst | BR   | sgn  | upd  | rsrv | 32b  | sgn  | rc    | lk   | priv | sgl  | rpt    |
|-----------------|------|--------|----------|------------------|-------------|-------------------|-------------|------|------|------|------|-------|------|-------|------|------|------|------|------|------|------|-------|------|------|------|--------|
|                 |      |        | op       |                  |             |                   |             |      | in   | out  | A    | out   | in   | out   | len  |      | ext  |      |      |      |      |       |      |      | pipe |        |
| INSN_illegal    | =>   | (ALU,  | NONE,    | OP_ILLEGAL,      | NONE,       | IMM, NONE,        | NONE, NONE, | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |
| INSN_fetch_fail | =>   | (LDST, | NONE,    | OP_FETCH_FAILED, | CIA,        | IMM, NONE,        | NONE, NONE, | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |
| INSN_add        | =>   | (ALU,  | NONE,    | OP_ADD,          | RA,         | RB, NONE,         | NONE, RT,   | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | RCOE, | '0', | '0', | '0', | NONE), |
| INSN_addc       | =>   | (ALU,  | NONE,    | OP_ADD,          | RA,         | RB, NONE,         | NONE, RT,   | '0', | '0', | '0', | '0', | ZERO, | '1', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | RCOE, | '0', | '0', | '0', | NONE), |
| INSN_adde       | =>   | (ALU,  | NONE,    | OP_ADD,          | RA,         | RB, NONE,         | NONE, RT,   | '0', | '0', | '0', | '0', | CA,   | '1', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | RCOE, | '0', | '0', | '0', | NONE), |
| INSN_addex      | =>   | (ALU,  | NONE,    | OP_ADD,          | RA,         | RB, NONE,         | NONE, RT,   | '0', | '0', | '0', | '0', | OV,   | '1', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | RC,   | '0', | '0', | '0', | NONE), |
| INSN_addg6s     | =>   | (ALU,  | NONE,    | OP_ADDG6S,       | RA,         | RB, NONE,         | NONE, RT,   | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |
| INSN_addi       | =>   | (ALU,  | NONE,    | OP_ADD,          | RA_OR_ZERO, | IMM, CONST_SI,    | NONE, RT,   | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |
| INSN_addic      | =>   | (ALU,  | NONE,    | OP_ADD,          | RA,         | IMM, CONST_SI,    | NONE, RT,   | '0', | '0', | '0', | '0', | ZERO, | '1', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |
| INSN_addic_dot  | =>   | (ALU,  | NONE,    | OP_ADD,          | RA,         | IMM, CONST_SI,    | NONE, RT,   | '0', | '0', | '0', | '0', | ZERO, | '1', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | ONE,  | '0', | '0', | '0', | NONE), |
| INSN_addis      | =>   | (ALU,  | NONE,    | OP_ADD,          | RA_OR_ZERO, | IMM, CONST_SI_HI, | NONE, RT,   | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |
| INSN_addme      | =>   | (ALU,  | NONE,    | OP_ADD,          | RA,         | IMM, CONST_M1,    | NONE, RT,   | '0', | '0', | '0', | '0', | CA,   | '1', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | RCOE, | '0', | '0', | '0', | NONE), |
| INSN_addpcis    | =>   | (ALU,  | NONE,    | OP_ADD,          | CIA,        | IMM, CONST_DXHI4, | NONE, RT,   | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |
| INSN_addze      | =>   | (ALU,  | NONE,    | OP_ADD,          | RA,         | IMM, NONE,        | NONE, RT,   | '0', | '0', | '0', | '0', | CA,   | '1', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | RCOE, | '0', | '0', | '0', | NONE), |
| INSN_and        | =>   | (ALU,  | NONE,    | OP_LOGIC,        | NONE,       | RB, NONE,         | RS, RA,     | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | RC,   | '0', | '0', | '0', | NONE), |
| INSN_andc       | =>   | (ALU,  | NONE,    | OP_LOGIC,        | NONE,       | RB, NONE,         | RS, RA,     | '0', | '0', | '1', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | RC,   | '0', | '0', | '0', | NONE), |
| INSN_andi_dot   | =>   | (ALU,  | NONE,    | OP_LOGIC,        | NONE,       | IMM, CONST_UI,    | RS, RA,     | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | ONE,  | '0', | '0', | '0', | NONE), |
| INSN_andis_dot  | =>   | (ALU,  | NONE,    | OP_LOGIC,        | NONE,       | IMM, CONST_UI_HI, | RS, RA,     | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | ONE,  | '0', | '0', | '0', | NONE), |
| INSN_attn       | =>   | (ALU,  | NONE,    | OP_ATTN,         | NONE,       | IMM, NONE,        | NONE, NONE, | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '1', | '1', | NONE), |
| INSN_brel       | =>   | (ALU,  | NONE,    | OP_B,            | CIA,        | IMM, CONST_LI,    | NONE, NONE, | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '1', | '0', | '0', | NONE), |
| INSN_babs       | =>   | (ALU,  | NONE,    | OP_B,            | NONE,       | IMM, CONST_LI,    | NONE, NONE, | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '1', | '0', | '0', | NONE), |
| INSN_bcrel      | =>   | (ALU,  | NONE,    | OP_BC,           | CIA,        | IMM, CONST_BD,    | NONE, NONE, | '1', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '1', | '0', | '0', | NONE), |
| INSN_bcabs      | =>   | (ALU,  | NONE,    | OP_BC,           | NONE,       | IMM, CONST_BD,    | NONE, NONE, | '1', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '1', | '0', | '0', | NONE), |
| INSN_bcctr      | =>   | (ALU,  | NONE,    | OP_BCREG,        | NONE,       | IMM, NONE,        | NONE, NONE, | '1', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '1', | '0', | '0', | NONE), |
| INSN_bclr       | =>   | (ALU,  | NONE,    | OP_BCREG,        | NONE,       | IMM, NONE,        | NONE, NONE, | '1', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '1', | '0', | '0', | NONE), |
| INSN_bctar      | =>   | (ALU,  | NONE,    | OP_BCREG,        | NONE,       | IMM, NONE,        | NONE, NONE, | '1', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '1', | '0', | '0', | NONE), |
| INSN_bperm      | =>   | (ALU,  | NONE,    | OP_BPERM,        | NONE,       | RB, NONE,         | RS, RA,     | '0', | '0', | '0', | '0', | ZERO, | '0', | NONE, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |
| INSN_brh        | =>   | (ALU,  | NONE,    | OP_BREV,         | NONE,       | IMM, NONE,        | RS, RA,     | '0', | '0', | '0', | '0', | ZERO, | '0', | is2B, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |
| INSN_brw        | =>   | (ALU,  | NONE,    | OP_BREV,         | NONE,       | IMM, NONE,        | RS, RA,     | '0', | '0', | '0', | '0', | ZERO, | '0', | is4B, | '0', | '0', | '0', | '0', | '0', | '0', | '0', | NONE, | '0', | '0', | '0', | NONE), |



# Static Branch Predict (Decode1)

```
-- Branch predictor
-- Note bclr, bcctr and bctar not predicted as we have no
-- count cache or link stack.
br_offset := f_in.insn(25 downto 2);
case icode is
  when INSN_brel | INSN_babs =>
    -- Unconditional branches are always taken
    v.br_pred := '1';
  when INSN_bcrel =>
    -- Predict backward relative branches as taken, others as untaken
    v.br_pred := f_in.insn(15);
    br_offset(23 downto 14) := (others => '1');
  when others =>
end case;
```

## Branch Conditional B-form

|      |                     |             |
|------|---------------------|-------------|
| bc   | BO, BI, target_addr | (AA=0 LK=0) |
| bca  | BO, BI, target_addr | (AA=1 LK=0) |
| bcl  | BO, BI, target_addr | (AA=0 LK=1) |
| bcla | BO, BI, target_addr | (AA=1 LK=1) |

| 16 | BO | BI | BD | AA | LK |
|----|----|----|----|----|----|
| 0  | 6  | 11 | 16 | 30 | 31 |

```
if (64-bit mode)
  then M ← 0
  else M ← 32
if ¬BO2 then CTR ← CTR - 1
ctr_ok ← BO2 | ((CTRM:63 ≠ 0) ⊕ BO3)
cond_ok ← BO0 | (CRBI+32 ≡ BO1)
if ctr_ok & cond_ok then
  if AA then NIA ←iea EXTS(BD || 0b00)
  else      NIA ←iea CIA + EXTS(BD || 0b00)
if LK then LR ←iea CIA + 4
```



# Immediate Field Formatting (Decode2)

```
function decode_b_const (t : const_sel_t; insn_in : std_ulogic_vector(31 downto 0);
                        prefix : std_ulogic_vector(25 downto 0))
    return std_ulogic_vector is
    variable ret : std_ulogic_vector(63 downto 0);
begin
    case t is
        when CONST_UI =>
            ret := std_ulogic_vector(resize(unsigned(insn_ui(insn_in)), 64));
        when CONST_SI =>
            ret := std_ulogic_vector(resize(signed(insn_si(insn_in)), 64));
        when CONST_PSI =>
            ret := std_ulogic_vector(resize(signed(insn_prefix_si(prefix, insn_in)), 64));
        when CONST_SI_HI =>
            ret := std_ulogic_vector(resize(signed(insn_si(insn_in)) & x"0000", 64));
        when CONST_UI_HI =>
            ret := std_ulogic_vector(resize(unsigned(insn_si(insn_in)) & x"0000", 64));
        when CONST_LI =>
            ret := std_ulogic_vector(resize(signed(insn_li(insn_in)) & "00", 64));
        when CONST_BD =>
            ret := std_ulogic_vector(resize(signed(insn_bd(insn_in)) & "00", 64));
        when CONST_DS =>
            ret := std_ulogic_vector(resize(signed(insn_ds(insn_in)) & "00", 64));
        when CONST_DQ =>
            ret := std_ulogic_vector(resize(signed(insn_dq(insn_in)) & "0000", 64));
        when CONST_DXHI4 =>
            ret := std_ulogic_vector(resize(signed(insn_dx(insn_in)) & x"0004", 64));
        when CONST_M1 =>
            ret := x"FFFFFFFFFFFFFFFF";
        when CONST_SH =>
            ret := x"0000000000000000" & "00" & insn_in(1) & insn_in(15 downto 11);
        when CONST_SH32 =>
            ret := x"0000000000000000" & "000" & insn_in(15 downto 11);
        when CONST_DSX =>
            ret := 55x"7FFFFFFFFFFFFFFF" & insn_in(0) & insn_in(25 downto 21) & "000";
        when others =>
            ret := (others => '0');
    end case;

    return ret;
end;
```

- Many different ways in which a constant encoded in the instruction can turn into an operand.
- Decoder ROM in decode1 determines which variant applies.
- Decode2 calculates immediate values based on the instruction and the type of constant.

# Execute1 and execute2

- Executes all instructions except load/store and FPU instructions
  - 64-bit integer multiply and integer division are now done by FPU (if present)
- Submodules handle shift/rotate instructions, logical instructions, count-bits instructions, 32-bit multiplications
- “Main adder” performs add/subtract instructions
- Exception detection and interrupt generation
- Interlocks to ensure instructions complete in order
- Execute1 does most integer computations and forwards results to decode2
  - Count-bits, multiply, and some move from SPR instructions take 2 cycles to generate their result, so their results are not forwarded from execute1
- Execute2 does updating of SPRs, condition-code result generation for dot-form instructions, and forwarding of results.



# Loadstore1/2 and dcache

- Loadstore1 receives instruction from execute1
  - Evaluates whether the access is misaligned and crosses a doubleword boundary
  - Converts FP data to single-precision form for stfs[x] instruction
- Effective address sent to dcache in first cycle
- Dcache reads the cache RAM, cache tag, and TLB in the second cycle
  - Dcache is store-through; all stores are written to memory immediately
  - Cache tags and TLB looked up in parallel; both can have multiple ways
    - Matrix of tag comparators to decide on TLB and cache hits
  - State machine to handle stores, load misses and non-cacheable loads
  - Store data forwarded to later loads to avoid stalling load-hit-store cases
- For loads, formatting of result data occurs in third (writeback) cycle
  - Except conversion of single-precision data to DP form (lfs[x]) takes an extra cycle
- Loadstore 2 (and 3) has state machine to handle complex cases

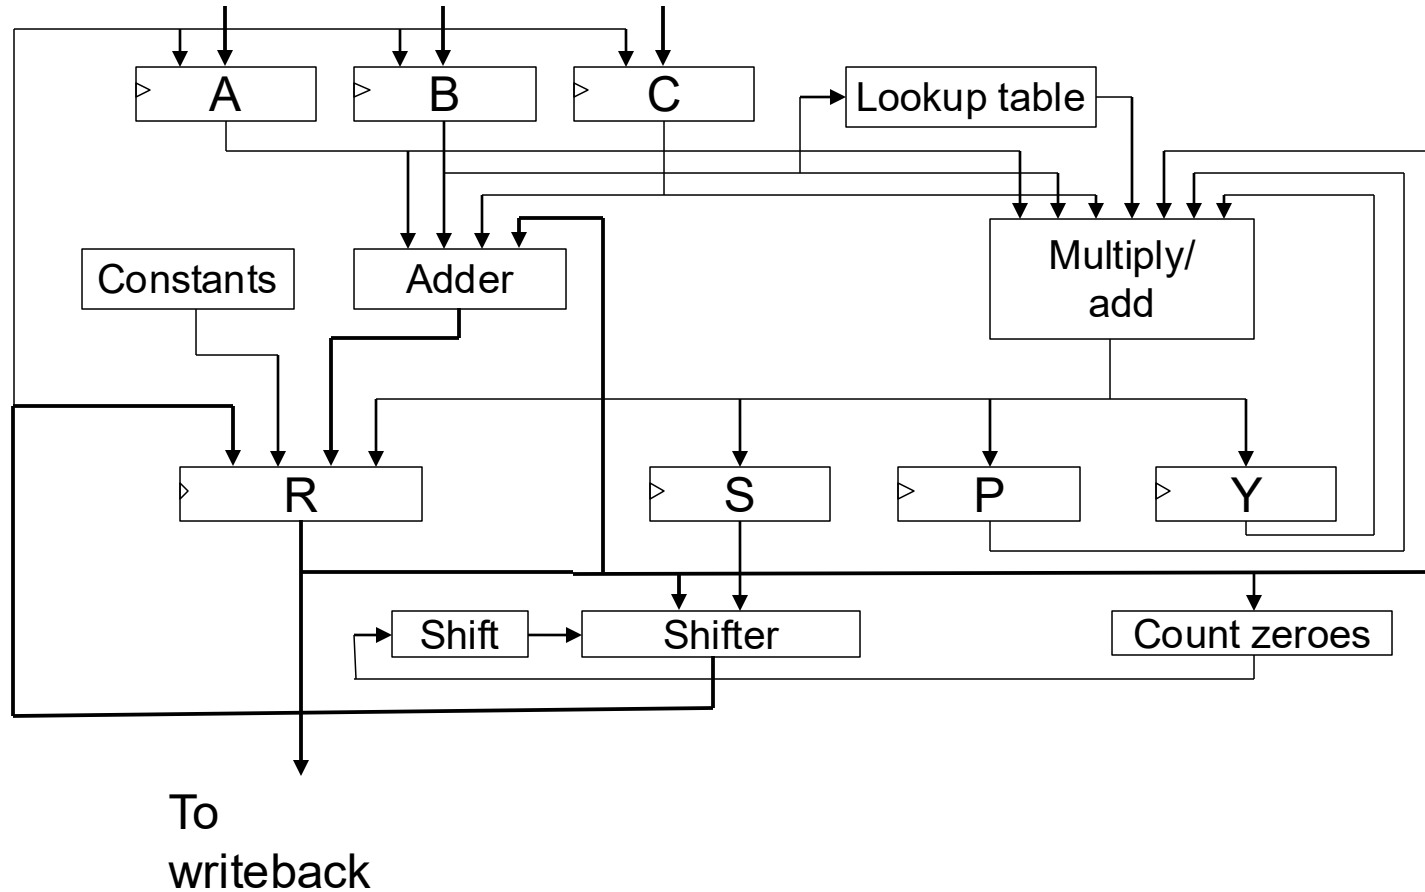


# Floating-point unit

- Non-pipelined
  - Has its own data path with 64-bit adder, shifter, multiplier ( $64 \times 64 \rightarrow 128$ ) and count-leading-zeroes logic
- Control is a state machine
  - Implemented with “random” logic (not microcode)
  - One operation at a time
  - Add/subtract generally take 5–12 cycles, multiply takes 8–15 cycles
- Handles all architected exception conditions
  - Denormalized operands and results
  - Underflow and overflow conditions, including trap-enabled and trap-disabled results



# FPU – Data paths (simplified)



# PMU

- Used for performance monitoring
- Receives information from/via the execute1 stage
- Maintains various counters dependent on MMCR settings
- Can raise an interrupt for certain events

```
type PMUToExecute1Type is record
    spr_val : std_ulogic_vector(63 downto 0);
    intr    : std_ulogic;
end record;

type Execute1ToPMUType is record
    mfspr : std_ulogic;
    mtspr : std_ulogic;
    spr_num : std_ulogic_vector(4 downto 0);
    spr_val : std_ulogic_vector(63 downto 0);
    tbbits : std_ulogic_vector(3 downto 0);
    pmm_msr : std_ulogic;
    pr_msr : std_ulogic;
    run : std_ulogic;
    nia : std_ulogic_vector(63 downto 0);
    addr : std_ulogic_vector(63 downto 0);
    addr_v : std_ulogic;
    trace : std_ulogic;
    occur : PMUEventType;
end record;
```

```
type PMUEventType is record
    no_instr_avail : std_ulogic;
    dispatch       : std_ulogic;
    ext_interrupt   : std_ulogic;
    instr_complete : std_ulogic;
    fp_complete    : std_ulogic;
    ld_complete    : std_ulogic;
    st_complete    : std_ulogic;
    br_taken_complete : std_ulogic;
    br_mispredict   : std_ulogic;
    ipref_discard   : std_ulogic;
    itlb_miss       : std_ulogic;
    itlb_miss_resolved : std_ulogic;
    icache_miss     : std_ulogic;
    dc_miss_resolved : std_ulogic;
    dc_load_miss    : std_ulogic;
    dc_ld_miss_resolved : std_ulogic;
    dc_store_miss   : std_ulogic;
    dtlb_miss       : std_ulogic;
    dtlb_miss_resolved : std_ulogic;
    ld_miss_nocache : std_ulogic;
    ld_fill_nocache : std_ulogic;
end record;
```




# Memory management unit

- Level-1 TLBs in fetch1 and dcache
  - 4kB page size, looked up in parallel with cache tags
  - Instruction TLB is direct-mapped, 64 entries
  - Data TLB is 2-way set associative, 128 entries total
- MMU is a state machine
  - Sends a series of requests to the dcache to read the process table and PTEs
  - Eventually sends the translation to the dcache or fetch1
    - EA bits inserted if page size > 4kB
- Currently no caching of PTEs or PDEs in the MMU
  - Nor any caching of partial translation results (no “page walk cache”)
  - Just the L1 iTLB and dTLB inside the instruction and data caches
- Microwatt implements vestigial (1 entry) partition table



# Memory management unit

- ISA defines two address translation schemes
  - Hashed page table (HPT), used by AIX and IBM/i and older Linux kernels
  - Radix page table, used by recent Linux kernels
- Microwatt implements radix and not HPT
- ISA defines a very general tree structure for radix trees





# Memory management unit

- Radix tree page directory entry format:

|   |   |  |     |  |     |
|---|---|--|-----|--|-----|
| V | L |  | NLB |  | NLS |
|---|---|--|-----|--|-----|

NLS = Next level size, number of bits used to index next level of tree

- Two standard layouts defined in ISA
  - 52-bit EA, 4 levels, 64k page size (index fields 13, 9, 9, 5 bits)
  - 52-bit EA, 4 levels, 4k page size (index fields 13, 9, 9, 9 bits)
- Microwatt implements a general radix tree walker state machine
  - Any NLS value from 5 to 16 is permitted
  - Address space size from 31 bits (2 GiB) to 62 bits (4 EiB)
  - Any power-of-2 page size  $\geq 4$  kiB
  - Any number of levels between 1 and 10

